

COMP353 Databases

Relational Algebra (RA) for Relational Data Model

1

Relational Algebra (RA)

- *Database Query languages* are specialized languages to ask for information (**queries**) in DB.
- **Relational Algebra (RA)** is a query language associated with the relational data model.
- Queries in RA are expressions using a collection of operators on relations in the DB.
- The input(s) and output of a RA query are relations
- A query is **evaluated** using the **current instance** of the input relations to produce the output

2

Operations in “standard” RA

- The well-known **set operations**
 - \cup **Union** (\cup)
 - \cap **Intersection** (\cap)
 - $-$ **Difference** ($-$)
 - **Special DB operations** that select “parts” of a relation instance
 - **Selection** (σ) – selects some rows (tuples) & discards the rest
 - **Projection** (π) – selects some columns (attributes) & discards the rest
 - Operations that “combine” the tuples from the argument relations
 - \times **Cartesian product** (\times) – pairs the tuples in all possible ways
 - **Join** (\bowtie) – pairs particular tuples from the two input relations
 - A unary operation to **rename** relations, called **Rename** (ρ)
- Note:** The output of a RA expression is an “unnamed” relation/set, i.e., RA expressions return **sets**, whereas **SQL** returns **multisets (bags)**

3

Compatibility Requirement

- We can apply the **set operators** of **union**, **intersection**, and **difference** to instances of relations **R** and **S** if **R** and **S** are *compatible*, that is they have “the same” schemas.
- **Definition:** Relations $S(A_1, \dots, A_n)$ and $R(B_1, \dots, B_m)$ are compatible if:
 - (1) $n=m$ and
 - (2) $\text{type}(A_i) = \text{type}(B_i)$ (or compatible types), for all $1 \leq i \leq n$.

4

Set Operations on Relations

Let **R** and **S** be relation schemas, and **r** and **s** be any instances of them.

- The **union** of **r** and **s** is the set of all tuples that appear in either one or both. Each tuple **t** appears only once in the union, even if it appears in both; $r \cup s = \{t \mid t \in r \vee t \in s\}$
- The **intersection** of **r** and **s**, is the set of all tuples that appear in both; $r \cap s = \{t \mid t \in r \wedge t \in s\}$
- The **difference** of **r** and **s**, is the set of all tuples that appear in **r** but not in **s**; $r - s = \{t \mid t \in r \wedge t \notin s\}$
 - **Commutative** operations; $r \cap s = s \cap r$
Note: Set difference ($-$) is not commutative, i.e., $(r - s) \neq (s - r)$

5

Example

Relation Schema: **Star** (name, address, gender, birthdate)

**Instance r
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

**Instance s
of Star:**

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

$r \cup s$:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88
Harrison Ford	789 Palm rd.	M	7/7/77

6

Example

Relation Schema: **Star** (name, address, gender, birthdate)

Instance **r**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

Instance **s**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

r ∩ s:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99

7

Example

Relation Schema: **Star** (name, address, gender, birthdate)

Instance **r**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

Instance **s**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

r - s:

Name	Address	Gender	Birthdate
Mark Hamill	456 Oak rd.	M	8/8/88

8

Example

Relation Schema: **Star** (name, address, gender, birthdate)

Instance **r**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Mark Hamill	456 Oak rd.	M	8/8/88

Instance **s**
of Star:

Name	Address	Gender	Birthdate
Carrie Fisher	123 Maple	F	9/9/99
Harrison Ford	789 Palm rd.	M	7/7/77

s - r:

Name	Address	Gender	Birthdate
Harrison Ford	789 Palm rd.	M	7/7/77

9

Projection (π)

- Let **R** be a relation schema.
- The **projection** operation (π) is used to produce, from any instance **r** of **R**, a new relation that includes listed "columns" of **R**
- The output of $\pi_{A_1, A_2, \dots, A_j}(r)$ is a relation with columns A_1, A_2, \dots, A_j , in this order.
- Note: The subscript of π is a *list*, which defines the structure of the output as the ordered tuple (A_1, A_2, \dots, A_j) .

10

Example

Relation Schema: **Movie** (title, year, length, filmType, studioName, producer)

Instance **movie**
Of Movie:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\pi_{\text{title, year, length}}(\text{movie})$

title	year	length
Star wars	1977	124
Mighty Ducks	1991	104
Wayne's World	1992	95

11

Example

Relation Schema: **Movie** (title, year, length, filmType, studioName, producer)

Instance **movie**
Of Movie:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\pi_{\text{filmType}}(\text{movie})$

Result:

filmType
color

12

Selection (σ)

- The **selection** operator (σ), applied to an instance r of relation R , returns a subset of r
- We denote this operation/query by $\sigma_C(r)$
- The output includes tuples satisfying condition C
- The schema of the output is the same as R

13

Example

Relation Schema: **Movie**(title, year, length, filmType, studioName, producer)

Instance **movie** of **Movie**:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\sigma_{\text{length} \geq 100}(\text{movie})$

Result:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890

14

Example

Relation: **Movie**(title, year, length, filmType, studioName, producer)

Instance **movie** of **Movie**:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

Query: $\sigma_{\text{length} \geq 100 \text{ AND studioName} = \text{'Fox'}}(\text{movie})$

Result:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345

15

Cartesian Product (\times)

- Let R and S be relation schemas, and r and s be any instances of R and S , respectively.
- The **Cartesian Product** of r and s is the set of all tuples obtained by “concatenating” the tuples in r and s . Formally, $r \times s = \{ t_1.t_2 \mid t_1 \in r \wedge t_2 \in s \}$
- The schema of result is the “union” of R and S
 - If R and S have some attributes in common, we need to invent new names for identical names, e.g., use $R.B$ and $S.B$, if B appears in both R and S

16

Example

Instance r of R :

A	B
1	2
3	4

Instance s of S :

B	C	D
2	5	6
4	7	8
9	10	11

$r \times s$:

A	R.B	S.B	C	D
1	2	2	5	6
1	2	4	7	8
1	2	9	10	11
3	4	2	5	6
3	4	4	7	8
3	4	9	10	11

17

Theta-join (θ)

- Suppose R and S are relation schemas, r is an instance of R , and s is an instance of S . The **theta-join** of r and s is the set of all tuples obtained from concatenating all $t_1 \in r$ and $t_2 \in s$, such that t_1 and t_2 satisfy some condition C
- We denote θ -join by $r \bowtie_C s$
- The schema of the result is the same as the schema of $R \times S$ (i.e., the union of R and S)
- C is a Boolean expression, simple or complex, as in operation σ

18

Example

Instance r of R :

A	B	C
1	2	3
6	5	8
9	7	11

Instance s of S :

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie_{A=D} s$:

A	R.B	R.C	S.B	S.C	D
1	2	3	2	3	4
1	2	3	2	3	5
1	2	3	7	8	10
6	5	8	7	8	10
9	7	11	7	8	10

19

Equi-join

- The **equi-join** operator, is a special case of θ -join, in which we may only use the equality relation ($=$) in condition C
- It is denoted as $r \bowtie_c s$ (i.e., the same as θ -join)
- The schema of the output is the same as that of θ -join

20

Example

Instance r of R :

A	B	C
1	2	3
6	5	8
9	7	11

Instance s of S :

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie_{R.C=S.C} s$:

A	R.B	R.C	S.B	S.C	D
1	2	3	2	3	4
1	2	3	2	3	5
6	5	8	7	8	10

21

Natural Join (\bowtie)

- **Natural join**, is a special case of equi-join, where the equalities are not explicitly specified, rather they are assumed implicitly on the common attributes of R and S
- We denote this natural join operation by $r \bowtie s$
- The schema of the output is similar to that of equi-join, except that each common attribute appears only once.

Note: If R and S do not have any common attribute, then the join operation becomes Cartesian product.

22

Example

Instance r of R :

A	B	C
1	2	3
6	2	8
9	7	3

Instance s of S :

B	C	D
2	3	4
2	3	5
7	8	10

$r \bowtie s$:

A	B	C	D
1	2	3	4
1	2	3	5

23

Expressing Queries in RA

- Every standard RA operation has relation(s) as argument(s) and produces a relation (set) as the output
(Exception is the **sort operator τ**)
- This property of RA operations (that inputs and outputs are relations) makes it possible to formulate/express any query by composing/nesting/grouping subqueries.
- We can use parentheses for *grouping*, in order to improve *clarity* and *readability*

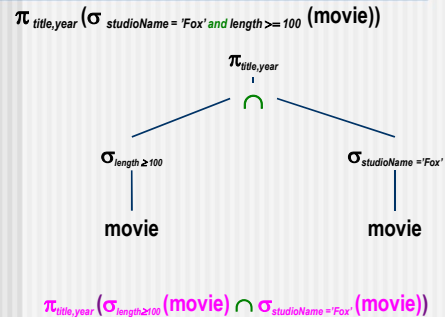
24

Example: RA Query

- Relation schema:
Movie (title, year, length, filmType, studioName)
- Query: List the title and year of every movie made by Fox whose length is at least 100 minutes?
- One way to express this query in RA is:
 $\pi_{\text{title, year}}(\sigma_{\text{studioName} = \text{'Fox' and length} \geq 100}(\text{movie}))$
- Another way:
 - Select those movie tuples that have length ≥ 100
 - Select those movie tuples that have studioName = 'Fox'
 - Find the intersection of the above two results
 - Then project on the attributes title and year

25

Example: RA Query



26

Example: RA Query

- Relation schema:
Movie (title, year, length, filmType, studioName)
StarsIn (title, year, starName)
- Query: List the names of the stars of movies of length ≥ 100 minutes long.
- One expression in RA for this query:
 - Select movie tuples of length ≥ 100
 - Join the result with relation StarsIn
 - Project on the attribute starName
- Exp1: $\pi_{\text{starName}}(\sigma_{\text{length} \geq 100}(\text{movie}) \bowtie \text{starsIn})$
- Another solution: $\pi_{\text{starName}}(\sigma_{\text{length} \geq 100}(\text{movie} \bowtie \text{starsIn}))$

27

Renaming Operator (ρ)

- To control manipulating the names of the attributes in formulating queries in relational algebra, we may need renaming of relations. May do this for convenience too
 - The Renaming Operator is denoted by $\rho_{\text{s}(A_1, A_2, \dots, A_n)}(r)$
 - The result is a copy of the input relation instance r , but renamed to s and its attributes to A_1, \dots, A_n , in that order.
 - Use $\rho_{\text{s}}(r)$ to give relation r a new name s (with the same attributes in r)
- That is, in this case, schema of s is the same as that of r .

28

Example

- Query: $\pi_{\text{starName}}(\sigma_{\text{length} \geq 100}(\text{movie}) \bowtie \text{starsIn})$
- This query can be rewritten in 2 steps as follows:
 - $\rho_{\text{M}}(\text{title, year, length, filmType, studioName})(\sigma_{\text{length} \geq 100}(\text{movie}))$
or even simpler as: $\rho_{\text{M}}(\sigma_{\text{length} \geq 100}(\text{movie}))$ if used in the same formula
 - Or use $M := \sigma_{\text{length} \geq 100}(\text{movie})$ as a separate formula and then formulate the query as: $\pi_{\text{starName}}(M \bowtie \text{starsIn})$
- Consider takes(sid, cid, grade)
- Query: Find ID of every student who has taken at least 2 courses.
- $\pi_{\text{takes.sid}}(\sigma_{(\text{takes.sid} = \text{T.sid and (takes.cid} \neq \text{T.cid) (takes} \times \rho_{\text{T}}(\text{takes}))})$

29

Dependent and Independent Operations

- Some RA operations can be expressed based on other operations. Examples include:
 - $r \cap s = r - (r - s)$
 - $r \bowtie_{\text{C}} s = \sigma_{\text{C}}(r \times s)$
 - $r \bowtie_{\text{L}} s = \pi_{\text{L}}(\sigma_{r.A_1 = s.A_1 \text{ AND } \dots \text{ AND } r.A_n = s.A_n}(r \times s))$,
where L is the list of attributes in R followed by those attributes in S that are not in R , and A_1, \dots, A_n are the common attributes of R and S

30

Relational Algebra with Bag Semantics

- Relations stored in DB are called **base** relations/tables.
- Base relations are normally sets; no duplicates.
- In some situations, e.g., during query processing, it is allowed for relations to have duplicate tuples.
- If duplicates are allowed in a collection, it is called **bag/multiset**.

Instance
r of R:

A	B	C
1	2	3
6	5	8
6	5	8
1	2	3
9	7	11

Here, r is a bag

31

Why Bags?

1. Faster projection operations

- Bag projection is faster, since otherwise returning distinct values is expensive (as we need sorting for duplicate elimination).

Another example: Computing the bag union $(r \cup^b s)$ is much cheaper than computing the standard set union $r \cup s$.

Formally, if r and s have n and m tuples, then the bag and set union operations will cost $O(n+m)$ and $O(n \cdot m)$, respectively.

2. Correct computation with some aggregation

- For example, to compute the *average* of values for attribute A in the previous relation, we must consider the bag of those values

32

Set Operations on Bags

- $r \cup^b s$, the **bag union** of r and s , is the bag of tuples that are in r , in s , or in both. If a tuple t appears n times in r , and m times in s , then t appears $n+m$ times in bag $r \cup^b s$

$$r \cup^b s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = n+m \}$$
- $r \cap^b s$, the **bag intersection** of r and s , is the bag of tuples that appear in both r and s . If a tuple t appears n times in r , and m times in s , then the number of occurrences of t in bag $r \cap^b s$ is $\min(n,m)$

$$r \cap^b s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \min(n,m) \}$$
- $r -^b s$, the **bag difference** of r and s is defined as follows:

$$r -^b s = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \max(0, n-m) \}$$

$$s -^b r = \{ t:k \mid t:n \in r \wedge t:m \in s \wedge k = \max(0, m-n) \}$$

33

Example

Bag r:	A	B	Bag s:	A	B	$r \cup^b s$:	A	B
	1	2		1	2		1	2
	3	4		3	4		3	4
	1	2		3	4		1	2
	1	2		5	6		1	2
							3	4
							3	4
							5	6

34

Example

Bag r:	A	B	Bag s:	A	B	$r \cap^b s$:	A	B
	1	2		1	2		1	2
	3	4		3	4		3	4
	1	2		3	4			
	1	2		5	6			

35

Example

Bag r:	A	B	Bag s:	A	B	$r -^b s$:	A	B
	1	2		1	2		1	2
	3	4		3	4		1	2
	1	2		3	4			
	1	2		5	6			

36

Example

Bag r:

A	B
1	2
3	4
1	2
1	2

Bag s:

A	B
1	2
3	4
3	4
5	6

$S \rightarrow r$:

A	B
3	4
5	6

37

Bag Projection π^B

- Let R be a relation scheme, and r be a collection of tuples over R , which could have duplicates.
- The **bag projection** operator is used to produce, from r , a bag of tuples over some of R .
- Even when r does not have duplicates, we may get duplicates when projecting on some attributes of R .
- That is, π^B does not eliminate the duplicates and hence corresponds exactly to the SELECT clause in SQL.

38

Example

Bag r:

A	B	C
1	2	5
3	4	6
1	2	7
1	2	8

$\pi_{A,B}^B(r)$:

A	B
1	2
3	4
1	2
1	2

39

Example

Relation Schema: **movie**(title, year, length, filmType, studioName, producer)

Instance:

title	year	length	filmType	studioName	producer
Star wars	1977	124	color	Fox	12345
Mighty Ducks	1991	104	color	Disney	67890
Wayne's World	1992	95	color	Paramount	99999

$\pi_{\text{filmType}}^B(\text{movie})$:

filmType
color
color
color

40

Selection on Bags

- The selection operator σ_C applied to an instance r of relation R , will return a subset of r
- The tuples in the output relation are those that satisfy condition C , which involves attributes of R
- Duplicates are **not** eliminated from the result of a bag-selection

Note: The selection operation σ in RA is different from the SELECT clause in SQL

41

Example

Bag r:

A	B	C
1	2	5
3	4	6
1	2	7
1	2	7

$\sigma_{C \geq 6}(r)$:

A	B	C
3	4	6
1	2	7
1	2	7

42

Cartesian Product of Bags

- The **Cartesian Product** of bags r and s is the bag of tuples that can be formed by concatenating pairs of tuples, the first of which comes from r and the second from s . In symbols, $r \times s = \{ t_1, t_2 \mid t_1 \in r \wedge t_2 \in s \}$
- Each tuple of one relation is paired with each tuple of the other, regardless of whether it is a duplicate or not
- If a tuple t_1 appears m times in a relation r , and a tuple t_2 appears n times in relation s , then tuple t_1, t_2 appears $m \cdot n$ times in their bag-product, $r \times s$

43

Example

Bag r :

A	B
1	2
1	2

Bag s :

B	C
2	3
4	5
4	5

$r \times s$:

A	R.B	S.B	C
1	2	2	3
1	2	2	3
1	2	4	5
1	2	4	5
1	2	4	5
1	2	4	5

44

Join of Bags

- The bag join is computed in the same way as the standard join operation
- Duplicates are not eliminated in a bag join operation

Bag r :

A	B
1	2
1	2

Bag s :

B	C
2	3
4	5

$r \bowtie s$:

A	B	C
1	2	3
1	2	3

45

Constraints on Relations

- RA offers a convenient way to express a wide variety of constraints, e.g., referential integrity and FD's.
- There are **two ways** to express constraints in RA
 1. If r is an expression in RA, then the constraint $r = \emptyset$ says:

" r has no tuples, i.e., or r is empty"
 2. If r and s are RA expressions, then the constraint $r \subseteq s$ says:

"every tuple in (the result of) r is in (the result of) s "

These constraints hold also when r and s are bags.

46

Constraints on Relations

- Note that these two types of constraints are not independent. Why?
- The constraint $r \subseteq s$ could also be written as $r - s = \emptyset$

This follows from the definition of " $-$ ", because $r \subseteq s$ iff $r - s = \emptyset$, meaning that there is no tuple in r which is not in s

47

Referential Integrity Constraints

- Referential integrity in relational data model means:
 - if there is a value v in a tuple t in a relation r , then it is expected that v appears in a particular component (attribute) of some tuple s in relation s

E.g., if tuple (s, c, g) is in table **takes**(*sid*, *cid*, *grade*), then there must be a **student** with *sid* = s and a **course** with *cid* = c such that s has taken c

IOW, the mentions of values s and c in **takes** "refers" to some values outside this relation, and these values **must** exist

48

Example

- Relation schemas:
Movie (title, year, length, filmType)
StarsIn (title, year, starName)
- Constraint:
the title and year of every movie that appears in relation starsIn must appear also in movie; otherwise there is a violation in referencing in starsIn
- Query in RA:
 - $\pi_{\text{title, year}}(\text{starsIn}) \subseteq \pi_{\text{title, year}}(\text{movie})$
 - or equivalently
 - $\pi_{\text{title, year}}(\text{starsIn}) - \pi_{\text{title, year}}(\text{movie}) = \emptyset$

49

Functional Dependencies

- Any functional dependency $X \rightarrow Y$ can be expressed as an expression in RA
- Example:
 Consider the relation schema:
Star (name, address, gender, birthdate)
- How to express the FD: **name** \rightarrow **address** in RA?

50

Functional Dependencies

- Relation schema:
Star (name, address, birthdate)
- With the FD: **name** \rightarrow **address**
- The **idea** is that if we construct all pairs of **star** tuples, we **must not** find a pair that agree on **name** but disagree on **address**
- To "construct" the pairs in RA, we use **Cartesian product**, and to find pairs that violate this FD, we use **selection**
- We are then ready to express this FD by equating the result to \emptyset , as follows...

51

Example

Star:

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

$\rho_{S1(\text{name, address, birthdate})}(\text{star})$ $\rho_{S2(\text{name, address, birthdate})}(\text{star})$

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

Name	Address	Birthdate
Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77

52

Example

s1 x s2:

S1.Name	S1.Address	S1.Birthdate	S2.Name	S2.Address	S2.Birthdate
Carrie Fisher	123 Maple	9/9/99	Carrie Fisher	123 Maple	9/9/99
Carrie Fisher	123 Maple	9/9/99	Mark Hamill	456 Oak rd.	8/8/88
Carrie Fisher	123 Maple	9/9/99	Harrison Ford	789 Palm rd.	7/7/77
Mark Hamill	456 Oak rd.	8/8/88	Carrie Fisher	123 Maple	9/9/99
Mark Hamill	456 Oak rd.	8/8/88	Mark Hamill	456 Oak rd.	8/8/88
Mark Hamill	456 Oak rd.	8/8/88	Harrison Ford	789 Palm rd.	7/7/77
Harrison Ford	789 Palm rd.	7/7/77	Carrie Fisher	123 Maple	9/9/99
Harrison Ford	789 Palm rd.	7/7/77	Mark Hamill	456 Oak rd.	8/8/88
Harrison Ford	789 Palm rd.	7/7/77	Harrison Ford	789 Palm rd.	7/7/77

$\sigma_{S1.name=S2.name \text{ AND } S1.address \neq S2.address}(s1 \times s2) = \emptyset$

53

Functional Dependencies

- Relation schema:
Star (name, address, birthdate)
- With the FD: **name** \rightarrow **address**
- In RA:
 - $\sigma_{S1.name=S2.name \text{ AND } S1.address \neq S2.address}(\rho_{S1}(\text{star}) \times \rho_{S2}(\text{star})) = \emptyset$

54

Domain Constraints

- Relation schema:
Star (name, address, gender, birthdate)
- How to express the following constraint?
Valid values for **gender** are 'F' and 'M'
- In RA:
 - $\sigma_{\text{gender} \neq 'F' \text{ AND } \text{gender} \neq 'M'}(\text{star}) = \emptyset$
 - This is an example of **domain constraints**

55

Domain Constraints

- Relation schema:
Employee (eid, name, address, salary)
- How to express the constraint:
Maximum employee salaries is \$150,000
- In RA:
 - $\sigma_{\text{salary} > 150000}(\text{employee}) = \emptyset$

56

"For All" Queries (1)

- Given the database schema:
Student(Sid, Sname, Addr)
Course(Cid, Cname, Credits)
Enrolled(Sid, Cid)
- Consider the query:
"Find students enrolled in **all** the courses."
- A first attempt (below) fails!
 $\pi_{\text{sid}}(\text{Enrolled})$
- This RA query returns students enrolled in some courses.
- So, how to correctly express "For All" types of queries?

57

"For All" Queries (2)

- A *solution strategy* would be to:
 - First find the list of "all" students (**all guys**), from which we then subtract those who have not taken at least a course (**bad guys**)
 - Then "**good guys**", would be "all guys" from which we remove the "bad guys", i.e.,
Answer (Good guys) = All guys – Bad guys

58

"For All" Queries (3)

- Set of all students that we should consider:
 $\text{All Courses} \leftarrow \pi_{\text{Cid}}(\text{Course})$
 $\text{All Students} \leftarrow \pi_{\text{Sid}}(\text{Student})$
- Steps to find **students not enrolled in all courses**
 - Create all possible "student-course" pairs:
 $\text{All: SC-Pairs} \leftarrow \pi_{\text{Sid}}(\text{Student}) \times \pi_{\text{Cid}}(\text{Course})$
 - Consider all "actual" student-course pairs (take from **Enrolled**)
 - Using 1 & 2, we then find students not enrolled in all courses:
Bad: $\pi_{\text{Sid}}(\pi_{\text{Sid}}(\text{Student}) \times \pi_{\text{Cid}}(\text{Course}) - \text{Enrolled})$
- Answer:** $\text{All} - \text{Bad}$

59

The Division Operation (\div)

- The previous query can be conveniently and expressed in RA using the **division** operator \div
 - Divide **Enrolled** by $\pi_{\text{Cid}}(\text{Course})$
that is, $\text{Enrolled} \div \pi_{\text{Cid}}(\text{Course})$
 - Schema of the result is {Sid, Cid} – {Cid}
- $R \div S$ requires that the attributes of **S** to be a subset of attributes of **R**.
 - The schema of the output would be $R - S$

60

Example: Enrolled (student, sport)

Find students enrolled in all sports {Hockey, Football}.

Enrolled (Student, sport)

Jim	Hockey
Joe	Football
Jim	Football
Sue	Hockey

61

Example: Enrolled(student, sport)

$$\pi_{\text{student}}(\text{Enrolled}) \times \pi_{\text{sport}}(\text{Enrolled}) - \text{Enrolled} =$$

Jim	Hockey
Jim	Football
Joe	Hockey
Joe	Football
Sue	Hockey
Sue	Football

Jim	Hockey
Joe	Football
Jim	Football
Sue	Hockey

Joe	Hockey
Sue	Football

$$\pi_{\text{student}}(\text{Enrolled}) - \pi_{\text{student}}(\pi_{\text{student}}(\text{Enrolled}) \times \pi_{\text{sport}}(\text{Enrolled}) - \text{Enrolled})$$

Jim
Joe
Sue

Joe
Sue

$$= ? \rightarrow \text{Jim is the only student enrolled in all sports}$$

62

Another Example

- $r \div s = \pi_{R-S}(r) - \pi_{R-S}(\pi_{R-S}(r) \times s - r)$
- Given the DB schema:
 - Customer(cid, name)
 - Branch(bid, district)
 - Account(cid, bid)
- Query: "Find the names of all customers who have an account in *every* branch located in the Westmount area"
- Solution?
 - $\pi_{\text{name}}(\text{Customer} \bowtie \text{Account} \bowtie (\sigma_{\text{district} = \text{"Westmount"}}(\text{Branch})))$?
- No, this returns the names of all customers who have an account at some branch in Westmount, but not necessarily at every such branch.

63

Database:

Customer(cid, name), Branch(bid, district), Account(cid, bid)

- We can apply the division operator \div
 - Find all customer-branch pairs (cid, bid) for which customer (cid) has an account at branch (bid):

$$\pi_{\text{cid,bid}}(\text{customer} \bowtie \text{account})$$
 - Divide the above by all bid's of branches in Westmount

$$\pi_{\text{bid}}(\sigma_{\text{district} = \text{"Westmount"}}(\text{branch}))$$
 - $\pi_{\text{name}}(((\text{customer} \bowtie \text{account}) \div \pi_{\text{bid}}(\sigma_{\text{district} = \text{"Westmount"}}(\text{branch}))))$
- Thus, the division operation $r \div s$ is defined as:

$$r \div s = \pi_{R-S}(r) - \pi_{R-S}(\pi_{R-S}(r) \times s - r)$$

64