

COMP353 Databases

Logical Query Languages: Datalog

Logical Query Languages (Section 5.3)

- Motivation
 - Logical **if-then** rules extend rather “naturally” and easily to **recursive queries**; Relational algebra doesn't!
 - Recursion is considered in SQL3
 - Logical rules (Datalog) form a basis for development of many concepts and techniques in database and knowledge base systems, with many applications such as **data integration**

Datalog

`AlongMovie(Title, Year) ← movie(Title, Year, Length, Type), Length ≥ 100.`

- The **head** – the left hand side of the arrow/implication
- The **body** – the right hand side is a conjunction (AND) of predicates (called subgoals)
- NOTE: The book uses **AND** in the rule bodies instead of commas.
`AlongMovie(Title, Year) ← movie(Title, Year, Length, Type) AND Length ≥ 100.`
- The **head** is a positive predicate (atom) and the subgoals in the rule **body** are **Atoms**
- Atom - a formula of the form $p(T_1, \dots, T_n)$, where p is a predicate and T_i 's are **terms**
 - Predicate – **normal** (ordinary) relation name (e.g., movie, p) or **built-in** predicates (e.g., \geq in the above example)
 - Terms (arguments) – In Datalog, T_i is either a **variable** or a **constant**
 - Subgoals in the rule body may be “negated” using NOT

Datalog

`longMovie(Title, Year) ← movie(Title, Year, Length, Type), Length ≥ 100.`

- A variable in a rule body is called **local** if it appears only in the rule body, e.g., **Length** and **Type**
- The head is “true” if there are values for local variables that make every subgoal (in the rule body) true
- If the body includes no negation, then the rule can be viewed as a join of relations in the rule body followed by a **projection** on the head variable(s)

Datalog

`longMovie(Title, Year) ← movie(Title, Year, Length, Type), Length ≥ 100.`

- This rule may be expressed in RA as:

$\rho_{\text{longMovie}}(\pi_{\text{Title, Year}}(\sigma_{\text{Length} \geq 100}(\text{movie})))$

Variable-Based Interpretations of Rules

- In principle, given the rule
- $r: H \leftarrow B_1, \dots, B_k$.
we consider all possible assignments I of values (constants in the domain) to the variables in the rule.
- Such assignments I are called **interpretations**.
- For every interpretation I of the rule, if the body is true under I , we add to the head relation, the tuple defined by H under I .
(we only consider **ground** interpretations/substitutions).
 - That is, if $I(B_i)$ is true, $\forall i \in \{1, \dots, k\}$, then $I(H)$ is true.
 - In this case, we say that “ I satisfies r ” or “ I is a model for r ”, (this is denoted as $I \models r$)

Example

$s(X, Y) \leftarrow r(X, Z), r(Z, Y), \text{NOT } r(X, Y).$

- Instance r:**
- | A | B |
|---|---|
| 1 | 2 |
| 2 | 3 |
- The only assignments that make the first subgoal true are:
 - $\mathcal{I}_1: X \rightarrow 1, Z \rightarrow 2$
 - $\mathcal{I}_2: X \rightarrow 2, Z \rightarrow 3.$
 - In case (1),
 - $Y \rightarrow 3$ makes the second subgoal $r(Z, Y)$ true
 - Since $(1, 3) \notin r$, then "**NOT** $r(X, Y)$ " is also true
 - Thus, we infer tuple $(1, 3)$ for the head relation, s
 - In case (2),
 - No value of "Y" makes the second subgoal true
- Instance s:**
- | A | B |
|---|---|
| 1 | 3 |

Tuple-Based Interpretations of Rules

- Consider tuple variables for each positive normal subgoals that range over their relations
- For each assignment of tuples to each of these subgoals, we determine the implied assignment θ of values to variables
- If the assignment \mathcal{I} is:
 - consistent and also
 - satisfies all the subgoals (normal and built-ins) in the body
 then we add to the head relation, the tuple defined by the head H under \mathcal{I}

Example

$s(X, Y) \leftarrow r(X, Z), r(Z, Y), \text{NOT } r(X, Y).$

- Instance r:**
- | A | B |
|---|---|
| 1 | 2 |
| 2 | 3 |
- Instance s:**
- | A | B |
|---|---|
| 1 | 3 |
- Have 4 assignments of tuples to subgoals:

$r(X, Z)$	$r(Z, Y)$
1. (1, 2)	(1, 2)
2. (1, 2)	(2, 3)
3. (2, 3)	(1, 2)
4. (2, 3)	(2, 3)
 - Only the second assignment
 - is consistent for the value assigned to Z and satisfies the negative subgoal "**NOT** $r(X, Y)$ "
 - $\rightarrow (1, 3)$ is the only tuple we get for s

Datalog Programs

- A datalog program is a finite collection of rules
- Note: while standard datalog does not allow negation, in our presentation here, the programs and rules are actually in datalog extended with negation and built-in predicates.
- Predicates/relations can be divided into two classes
 - EDB** Predicates (input relations), also called **FACTS**
 - Extensional database = relations stored explicitly in DB
 - IDB** Predicates (derived/output relations), defined by rule(s)
 - Intensional database
 - They are similar to views in relational databases
- Note:** EDB predicates appear only in the rule body and IDBs appear in the head and possibly in the body

Operations in Datalog

- The usual set operations
- Consider relation schemas $r(X, Y)$ and $s(X, Y)$
 - Intersection**
 - RA: $Q = r \cap s$
 - Datalog: $q(X, Y) \leftarrow r(X, Y), s(X, Y).$
 - Union**
 - RA: $Q = r \cup s$
 - Datalog: the following two rules:
 - $q(X, Y) \leftarrow r(X, Y).$
 - $q(X, Y) \leftarrow s(X, Y).$
 - Difference**
 - RA: $Q = r - s$
 - Datalog: $q(X, Y) \leftarrow r(X, Y), \text{NOT } s(X, Y).$

Operations in Datalog

- Projection operation
 - RA: $p = \pi_x(r)$
 - Datalog: $p(X) \leftarrow r(X, Y).$

Operations in Datalog

■ Selection operation

- RA: $s = \sigma_{X > 10 \text{ AND } Y = 5}(r)$
- Datalog: $s(X,Y) \leftarrow r(X,Y), X > 10, Y = 5.$

Operations in Datalog

■ Selection operation. Recall the schema of $r(X,Y)$.

- RA: $s = \sigma_{X > 10 \text{ OR } Y = 5}(r)$
- Datalog: the following two rules:
 1. $s(X,Y) \leftarrow r(X,Y), X > 10.$
 2. $s(X,Y) \leftarrow r(X,Y), Y = 5.$

Note: the following Datalog program is equivalent to the above.

1. $s(A,B) \leftarrow r(A,B), A > 10.$
2. $s(C,D) \leftarrow r(C,D), D = 5.$

Operations in Datalog

■ Cartesian Product operation

■ Consider relation schemas $r(A, B)$ and $s(C, D)$

- RA: $Q = r \times s$
- Datalog: $q(X, Y, Z, W) \leftarrow r(X,Y), s(Z,W).$

Operations in Datalog

■ Join operation

■ Theta-join with an **AND** condition, e.g., " c_1 AND c_2 "

- RA: $tj1 = r \bowtie_{X > Z \text{ AND } Y < W} s$
- Datalog: $tj1(X, Y, Z, W) \leftarrow r(X,Y), s(Z,W), X > Z, Y < W.$

■ Theta-join with an **OR** condition, e.g., " c_1 OR c_2 "

- RA: $tj2 = r \bowtie_{X > Z \text{ OR } Y < W} s$
- Datalog: the following two rules:
 - $tj2(X, Y, Z, W) \leftarrow r(X,Y), s(Z, W), X > Z.$
 - $tj2(X, Y, Z, W) \leftarrow r(X,Y), s(Z,W), Y < W.$

Operations in Datalog

■ Join operation

■ Equi-join

- RA: $ej3 = r \bowtie_{Y=Z} s$
 - Datalog: $ej3(X,Y,Z,W) \leftarrow r(X,Y), s(Z,W), Y = Z.$
- OR even better (simpler):
- $$ej3(X,Y,Y,W) \leftarrow r(X,Y), s(Y,W).$$

Operations in Datalog

■ Join operation

■ Natural join

- RA: $nj4 = r \bowtie s$
- Datalog: $nj4(X,Y,W) \leftarrow r(X,Y), s(Y,W).$

Example:Datalog Queries/Programs

- Database schema:
movie(Title, Year, Length, FilmType, StudioName)
starsIn(Title, Year, StarName)
 - Query: Find the names of stars of movies that are at least 100 minutes long
 - Relational Algebra Expression:
 $Q = \pi_{starName} (\sigma_{length \geq 100} (movie) \bowtie starsIn)$
 - Datalog program:
 $r1(Title, Year, Length, Type, Studio) \leftarrow movie(Title, Year, Length, Type, Studio), Length \geq 100.$
 $r2(Title, Year, Length, Type, Studio, Name) \leftarrow r1(Title, Year, Length, Type, Studio), starsIn(Title, Year, Name).$
 $q(Name) \leftarrow r2(Title, Year, Length, Type, Studio, Name).$
- As in RA case, we could express this query using just one rule, as follows:
 $q(Name) \leftarrow movie(Title, Year, Length, Type, Studio), Length \geq 100, starsIn(Title, Year, Name).$

Expressive Power of Datalog


- Relational algebra = Nonrecursive Datalog+ negation
- Datalog can express SQL **SELECT-FROM-WHERE** statements that do **not** use aggregation and/or grouping
- The SQL-99 standard supports recursion but it is not part of the "core" SQL-99 standard that every DBMS should support
- Some DBMS implementations, e.g. DB2, support *linear recursion*

Example

- trainSchedule(From,To) Find every cities that can be reached from Montreal by train
- Datalog:
 $fromMontreal(C) \leftarrow trainSchedule('Montreal', C).$
 $fromMontreal(TC) \leftarrow fromMontreal(C), trainSchedule(C,TC).$

trainSchedule:	From	To	fromMontreal:
	Toronto	Calgary	
	NYC	Boston	
	NYC	Albany	
	Chicago	Detroit	
	Montreal	Toronto	
	Montreal	NYC	
	Boston	Quebec City	

Evaluation of Recursive Rules



Example

Relation Schema: **sequelOf**(Movie,Sequel)

Instance:

Movie	Sequel
Star wars	Star wars II
Naked Gun	Naked Gun 2
Naked Gun 2	Naked Gun 2 ½
Star wars II	Star wars III
Naked Gun 2 ½	Naked Gun 3

For a given movie, find all the follow-up movies, i.e., a sequel, a sequel of a sequel, and so on

Example

Relation Schema: **sequelOf**(Movie, Sequel)

Instance:

Movie	Sequel
Star wars	Star wars II
Naked Gun	Naked Gun 2
Naked Gun 2	Naked Gun 2 ½
Star wars II	Star wars III
Naked Gun 2 ½	Naked Gun 3

$followUp(X, Y) \leftarrow sequelOf(X,Y).$
 $followUp(X, Y) \leftarrow sequelOf(X,Z), followUp(Z,Y).$

Recursion

- Let P be any datalog program
- We say an IDB predicate r in P *depends on* predicate s if there is a rule in P with r as the head and s as a subgoal in the rule body
- Construct the (dependency) graph of P :
 - **Nodes** -- IDB predicates in P
 - **Arcs** -- an arc from node r to s if r depends on s
 - Label the arc with ' \neg ' for negated subgoals
- P is recursive iff its dependency graph has a cycle

Example

```
followUp(X, Y) ← sequelOf(X, Y).  
followUp(X, Y) ← sequelOf(X, Z), followUp(Z, Y).
```



Safety

- It is possible to write a rule that makes "no sense".
- Example of such rules:
 - $s(X) \leftarrow r(Y)$.
 - $s(X) \leftarrow \text{NOT } r(X)$.
 - $s(X) \leftarrow r(Y), X < Y$.
- In each of these rules, the IDB relation s (output relation) could be infinite, even if (the input) relation r is finite
- Such rules are said to be not *SAFE*

Safety

- For a rule to be safe, the following conditions must hold:
 - If a variable X appears in the rule head, then X must appear in an "ordinary" predicate in the body or be equal to such a variable (directly or indirectly), e.g., $X=Y$, and Y appears in an ordinary predicate in the rule body.

Recall: the predicates could be ordinary or built-in.